

SafeContract: Composable Action-Space Monitors for VLA Deployment

Krishnam Gupta
Independent Research
San Francisco, CA
kayjee1994@gmail.com

Abstract—Vision-Language-Action (VLA) models predict robot actions from images and instructions, but no standardized action-space monitoring exists for their deployment. We introduce SafeContract, a design-by-contract [5] monitoring toolkit that wraps any VLA policy with composable action-space contracts via a Python decorator. We show that the monitoring signal from a lightweight action-space contract reveals task-specific behavioral patterns and detects distribution shift. Bounds enforcement is a byproduct; the primary value is observability. Controlled ablations demonstrate zero false positives on clean policies while catching all violations on drifting and jerky policies. Across 130 LIBERO tasks, each task produces a distinct violation fingerprint across joint dimensions, with bounds violation rates varying from 4% to 37% and velocity violations dominating at 72.4%. Cross-model experiments reveal architecture-specific signatures: ACT produces 0% open-loop violations but 3,949 in closed-loop from compounding errors, Diffusion policies show 100% mismatch violations, and SmolVLA’s 98% velocity rate is 6× higher than ground-truth (16%). The full decorator adds $\sim 13\mu\text{s}$ overhead, requiring no QP solver, VLM, or model retraining. We release the toolkit as open source.¹

Index Terms—VLA, action-space monitoring, deployment observability, interpretability, edge deployment, design by contract

I. Introduction

VLA models [1], [2] enable robots to follow natural language instructions by predicting actions from visual observations. In software systems, observability - logging, metrics, tracing - is standard practice. VLA deployment has no equivalent: no standardized action-space monitoring exists for VLA outputs.

This is not theoretical. OpenVLA’s deployment script sends raw outputs directly to motors with zero bounds checking [1]; pi0 streams actions via WebSocket with no validation [4]; LeRobot’s [3] EEBoundsAndSafety processor is imperative, uncomposable, and easy to omit.

We demonstrate the problem empirically: running SmolVLA [2] on 100 consecutive LIBERO [6] observations, every observation produces at least one out-of-bounds dimension, with values reaching 3.70 (Sec. IV-A).

The deeper problem is observability. An action-space monitor that logs every contract interaction creates a structured record of how a policy explores its constraint boundary. We show this record is rich: different tasks

produce distinct violation fingerprints across joints, and shifts in violation rate signal distribution change.

We introduce SafeContract with three contributions:

- 1) A composable action-space monitoring toolkit via a `@safety_contract` Python decorator on any VLA `predict()` method (Properties 1–3).
- 2) Bounds enforcement as a byproduct: zero false positives on clean policies, full violation detection on corrupted ones (Sec. IV-C).
- 3) Violation fingerprinting: 130 LIBERO tasks produce task-specific joint-violation signatures enabling distribution shift detection (Sec. IV-D).

II. Method

A. Safety Contracts

A safety contract $C = (\mathbf{l}, \mathbf{u}, v_{\max})$ specifies per-joint action bounds and a velocity limit: $l_i \leq a_i \leq u_i$ for each joint i , and $\|\mathbf{a}_t - \mathbf{a}_{t-1}\|_{\infty} \leq v_{\max}$.

B. Enforcement Order

Enforcement proceeds in three phases: (1) clip to per-joint bounds, (2) clamp velocity (project each joint’s delta onto $[-v_{\max}, v_{\max}]$), (3) re-clip to bounds. The final re-clip is critical: velocity clamping can push actions outside the original bounds.

The decorator pattern means any VLA adapter gains monitoring with one line:

```
@safety_contract(action_range=[-1,1],  
                  joint_velocity_max=0.1)  
def predict(self, image, instruction):  
    return self.model(image, instruction)
```

Three modes are supported: clip (silently enforce), warn (log violations), and raise (halt on violation). All violations are logged with timestep, dimension, magnitude, and severity.

C. Composition Properties

Property 1 (Composition via Intersection). For hyperrectangular contracts $C_1 = (\mathbf{l}_1, \mathbf{u}_1)$, $C_2 = (\mathbf{l}_2, \mathbf{u}_2)$, sequential clipping equals intersection clipping: $\text{clip}(\text{clip}(\mathbf{a}, C_1), C_2) = \text{clip}(\mathbf{a}, C_1 \cap C_2)$. Stacking `@safety_contract` decorators is order-independent for per-joint bounds; velocity composition requires merging into a single contract with $v_{\max} = \min_i(v_{\max,i})$.

¹Code: <https://github.com/krishnam94/vla-edge>

Proof. Follows from idempotence and commutativity of per-dimension min/max. Verified empirically on 200 random-walk sequences (1,482 violations). $\square$

Property 2 (Feasibility from Safe States). Let $\mathbf{a}_{t-1}$ satisfy contract C with $v_{\max} > 0$. The feasible set is non-empty, since the “hold position” action $\mathbf{a}_t = \mathbf{a}_{t-1}$ satisfies both bounds and velocity constraints.

Property 3 (Shift Detection Bound). For a stationary policy π and fixed contract C , let each step produce an independent Bernoulli violation indicator with true rate r^* . By Hoeffding’s inequality [20], the empirical rate $\hat{r}_w$ over w steps satisfies

$$\Pr(|\hat{r}_w - r^*| > \epsilon) \leq 2\exp(-2w\epsilon^2). \quad (1)$$

At confidence $1-\alpha = 0.95$: $w \geq 738$ for $\epsilon = 0.05$, or $w \geq 185$ for $\epsilon = 0.1$.

Caveat. VLA actions exhibit temporal autocorrelation, so the effective sample size may be smaller than w . For β -mixing processes the bound degrades by a factor depending on the mixing rate [21]. In practice, concentration holds well empirically (Sec. IV-D).

III. Related Work

SafeVLA [7] fine-tunes with constrained RL (83% cost reduction) and RobustVLA [10] adds robust optimization, but both require retraining. At inference time, AEGIS [8] solves a CBF-QP via VLM scene understanding (0.36 ms) - for box constraints this reduces to per-dimension clipping, identical to SafeContract at $\sim 28\times$ higher cost; SafeDec [9] applies STL-constrained decoding (requires logit access); Safety Chip [11] uses LTL monitors for plan-level constraints; SafeDiffuser [12] and CoDiG [13] embed barriers in diffusion denoising; ATACOM [14] projects onto constraint manifolds; and ABC [17] brings design-by-contract to discrete LLM decisions. For runtime monitoring, Sentinel [18] combines temporal consistency with VLM monitoring but does not prevent unsafe actions; LIBERO-Plus [15] and VLA-Arena [16] benchmark VLA robustness; SABER [19] demonstrates black-box adversarial attacks on VLAs.

SafeContract differs from all the above: (i) architecture-agnostic, (ii) explicit composition, (iii) pure clipping ($\sim 13\mu\text{s}$), (iv) violation logging as deployment observability. It is the only method that monitors and enforces without model internals, external sensors, or retraining.

IV. Experiments

Platform: Apple M3, CPU, Python 3.12. Full contract overhead: $15.6 \pm 3.8\mu\text{s}$ (5,000 iterations), 0.000075% of SmolVLA inference time.

A. VLAs Produce Unsafe Actions (EXP-A)

We run SmolVLA on 100 consecutive LIBERO observations. Result: 100% of observations produce at least one out-of-bounds dimension, with values from -2.16 to 3.70

SmolVLA outputs exceed safe bounds on LIBERO




Fig. 1: SmolVLA action ranges on 100 LIBERO observations. All 6 dimensions exceed $[-1, 1]$ bounds (green dashes), reflecting normalization mismatch.

TABLE I: Controlled ablation. Clean policy: zero false positives. Corrupted policies: fully caught. Naive np.clip misses all velocity violations.

Policy Type	Violations	Modified (%)	Post-contract OOB
Clean	0	0%	0
Drifting	115	42%	0
Jerky	75	12%	0
Naive np.clip	0 caught	–	199 vel. missed

(expected $[-1, 1]$). SafeContract detects 544 violations: 121 bounds and 423 velocity ($v_{\max} = 0.1$).

Normalization analysis. SmolVLA was trained on SO-100 data. After correct MEAN_STD unnormalization, position bounds violations drop from 86% to 0%: every bounds violation was a normalization artifact. However, velocity violations persist: 16% of expert demonstration steps and 98% of SmolVLA outputs exceed $v_{\max}=0.1$ on at least one joint. The $6\times$ gap quantifies how much jerkier the learned policy is than training data. Bounds enforcement catches integration bugs; velocity monitoring catches physical concerns normalization cannot address.

B. Closed-Loop Task Success (EXP-CL)

We evaluate the ACT policy (51.6M params) on ALOHA sim transfer-cube, 50 episodes per condition. Without SafeContract: 58% success (29/50). With SafeContract (data-driven bounds, $v_{\max}=0.05$): 60% success (30/50). Fisher’s exact test: $p = 1.0$. SafeContract caught 3,949 velocity violations and modified 3,891 actions while maintaining identical task completion (Fig. 2).

C. Controlled Ablation (EXP-G)

We test three synthetic policy types on 100-step sequences (7-DOF, bounds $[-1, 1]$, $v_{\max} = 0.1$):

The clean policy produces zero violations: a safe policy passes through untouched. The drifting policy (cumulative noise) produces 115 violations (42% corrected); the jerky

Closed-Loop: ACT on ALOHA Transfer Cube (n=50)




Fig. 2: Closed-loop evaluation: ACT on ALOHA sim (50 episodes per condition). SafeContract catches 3,949 violations with no task success degradation ($p=1.0$).




Fig. 3: Noise sensitivity. Violations scale monotonically with noise; SafeContract eliminates 100% at every level.

policy (random spikes) produces 75 velocity violations (12% smoothed). Post-contract: zero OOB in all cases.

Noise sensitivity. Gaussian noise at 7 levels ($\sigma \in \{0, \dots, 1.0\}$) on 500 ground-truth LIBERO actions: violations scale monotonically from 527 to 19,269, and SafeContract eliminates 100% at every level (Fig. 3).

D. Violation Fingerprinting (EXP-H)

Violation patterns are informative, not just corrective.

Synthetic tasks. Four manipulation tasks (reaching, stacking, drawer, pouring) run for 200 steps each with bounds $[-1, 1]$, $v_{\max} = 0.1$. Each produces a unique joint-violation fingerprint (Fig. 4): reaching shows broad shoulder/elbow OOB; stacking shows wrist velocity violations; drawer shows a single-joint spike; pouring shows isolated pitch violations. Fingerprints are stable across 5 seeds (CV < 0.15 for 3 of 4 tasks).

Shift detection. Non-overlapping windows of 20 steps yield $p < 0.05$ for five of six task pairs ($p < 0.002$ for three; drawer vs. pouring not significant at $p = 0.34$). Note: the binomial test here operates on 10 independent




Fig. 4: Normalized violation fingerprints. Left: bounds violations vanish after unnormalization. Right: velocity fingerprints persist and differ by task.




Fig. 5: Bounds violation rates across 90 LIBERO-90 tasks (sorted). Violation rates range from 4% (drawer manipulation) to 37% (mug placement), demonstrating task-dependent monitoring signals.

window-level violation rates (each aggregating 20 steps), not on per-step Bernoulli trials; the Hoeffding bound in Property 3 applies to per-step concentration and would require larger w for the same ϵ . Nearest-centroid classification achieves 100% accuracy across 40 trials.

130 LIBERO tasks. We ran SmolVLA on all 130 tasks (40 standard + 90 diverse). Bounds violation rates vary from 4% to 37%; velocity violations dominate at 72.4%. Rotational axes are most violated (pitch 30.6%, roll 26.6%) while lateral motions are least affected (yaw 9.3%, y 9.7%). After correct unnormalization, bounds violations drop to 0%, but velocity fingerprints persist by suite: spatial ($x=83\%$, $z=85\%$), object ($y=85\%$), goal ($z=78\%$), 10-task (gripper=46%). Cosine distance between suite fingerprints is 0.054, confirming task-dependent patterns (Fig. 5).

Cross-model comparison. The violation signature is architecture-dependent. ACT produces 0% open-loop violations (smooth trajectories) but 3,949 velocity violations in closed-loop, where compounding prediction errors amplify jerkiness. Diffusion policies show 0% normalized bounds violations but 100% mismatch violations and

12% velocity violations. SmolVLA shows 98% velocity violations, $6\times$ higher than ground-truth (16%). Each architecture produces a distinct profile that SafeContract surfaces without model internals: ACT’s violations emerge only in closed-loop, Diffusion policies produce scale mismatches, and SmolVLA produces persistent jerkiness.

V. Discussion

SafeContract provides deployment observability for VLA action spaces: one-line adoption, composable per-joint bounds (Property 1), $\sim 13\mu\text{s}$ overhead. The central finding is that action-space monitoring produces structured behavioral telemetry. Different tasks produce distinct violation fingerprints, and violation rate changes detect distribution shift.

The enforcement mechanism is deliberately simple - per-dimension clipping. The ablation confirms it adds no distortion to safe policies while catching all violations in corrupted ones. Enforcement is a byproduct; the value is in the monitoring signal.

Normalization vs. monitoring. Normalization eliminates most bounds violations. SafeContract is complementary: velocity enforcement addresses physical hardware constraints normalization cannot, and violation logging enables debugging and shift detection. Tighter contracts catch more violations but clip more aggressively; at the moderate operating point (± 1.0 , $v_{\max}=0.1$), SafeContract catches 2,019 violations while correcting by less than half a unit on average.

Behavioral telemetry. SafeContract’s violation logs constitute behavioral telemetry for VLA deployment, analogous to observability tooling in software, UEBA in cybersecurity, or statistical process control in manufacturing. Baseline rates characterize normal behavior; deviations signal distribution shift, model regression, or integration errors. The 4%-37% spread across 130 tasks, with velocity dominating at 72.4%, provides a deployment-time signature unavailable from training metrics. To our knowledge, no prior work treats VLA action-space constraint interactions as an observability signal.

Limitations. (1) Synthetic tasks in EXP-G/H; real patterns may differ. (2) Closed-loop on one architecture (ACT/ALOHA, $p=1.0$); SmolVLA closed-loop limited by checkpoint issues. (3) 100% OOB primarily reflects normalization mismatch. (4) Violation rate observation is empirical, not a formal guarantee.

Future work. PropertyGuard (property-based fuzzing against contracts), cross-architecture fingerprinting (SmolVLA vs. OpenVLA vs. pi0), conformal prediction calibration.

References

- [1] Kim, M.J. et al., “OpenVLA: An Open-Source Vision-Language-Action Model,” CoRL 2024, arXiv:2406.09246.
- [2] Shukor, M. et al., “SmolVLA: A Vision-Language-Action Model for Affordable and Efficient Robotics,” arXiv:2506.01844, 2025.
- [3] Cadene, R. et al., “LeRobot: State-of-the-art Machine Learning for Real-World Robotics in PyTorch,” 2024, <https://github.com/huggingface/lerobot>.
- [4] Black, K. et al., “ π_0 : A Vision-Language-Action Flow Model for General Robot Control,” arXiv:2410.24164, 2024.
- [5] Meyer, B., “Applying ‘Design by Contract’,” Computer 25(10):40–51, 1992.
- [6] Liu, B. et al., “LIBERO: Benchmarking Knowledge Transfer for Lifelong Robot Learning,” NeurIPS 2023.
- [7] Zhang, B. et al., “SafeVLA: Towards Safety Alignment of Vision-Language-Action Model via Constrained Learning,” NeurIPS 2025, arXiv:2503.03480.
- [8] Hu, S. et al., “VLSA: VLA Models with Plug-and-Play Safety Constraint Layer,” arXiv:2512.11891, 2025.
- [9] Kapoor, P. et al., “Constrained Decoding for Robotics Foundation Models,” arXiv:2509.01728, 2025.
- [10] Guo, J. et al., “On Robustness of Vision-Language-Action Models against Multi-Modal Perturbations,” ICLR 2026, arXiv:2510.00037.
- [11] Yang, Z. et al., “Plug in the Safety Chip: Enforcing Constraints for LLM-driven Robot Agents,” ICRA 2024, arXiv:2309.09919.
- [12] Xiao, W. et al., “SafeDiffuser: Safe Planning with Diffusion Probabilistic Models,” ICLR 2025, arXiv:2306.00148.
- [13] Ma, H. et al., “CoDiG: Constraint-Aware Diffusion Guidance for Robotics,” CoRL 2025, arXiv:2505.13131.
- [14] Tolle, M. et al., “Towards Safe Robot Foundation Models Using Inductive Biases,” arXiv:2505.10219, 2025.
- [15] Fei, S. et al., “LIBERO-Plus: In-depth Robustness Analysis of VLA Models,” arXiv:2510.13626, 2025.
- [16] Zhang, B. et al., “VLA-Arena: An Open-Source Framework for Benchmarking VLA Models,” arXiv:2512.22539, 2025.
- [17] Bhardwaj, V.P., “Agent Behavioral Contracts,” arXiv:2602.22302, 2026.
- [18] Agia, C. et al., “Unpacking Failure Modes of Generative Policies: Runtime Monitoring of Consistency and Progress,” CoRL 2024, arXiv:2410.04640.
- [19] Wu, X. et al., “SABER: A Stealthy Agentic Black-Box Attack Framework for VLAs,” arXiv:2603.24935, 2026.
- [20] Hoeffding, W., “Probability Inequalities for Sums of Bounded Random Variables,” J. Amer. Statist. Assoc. 58(301):13–30, 1963.
- [21] Yu, B., “Rates of Convergence for Empirical Processes of Stationary Mixing Sequences,” Ann. Probab. 22(1):94–116, 1994.